

Benchmarking LLM Safety Against Social Engineering Techniques

Project Report (Group)

Client: CSIRO Data61 Australia

Unit: COS80029 - Technology Application Project

Submitted by: Team XC2

Facilitator: Dr. Xinyi Cai

Date: 01/11/2025

The Team

Name	Student ID	Contact Email
Vaisakh Pariyacheri	104800822	104800822@student.swin.edu.au
Aarathi Narayanan Kadambatta	104801951	104801951@student.swin.edu.au
Vikas Kadruvan	104981442	104981442@student.swin.edu.au
Samsun Gulshan Sheik Dawood	105009251	105009251@student.swin.edu.au
Arun Ragavendhar Arunachalam Palaniyappan	104837257	104837257@student.swin.edu.au
Aayudh Himanshubhai Vaghani	104996006	104996006@student.swin.edu.au

Executive Summary

This report details an investigation into the safety robustness of commercial LLMs (Grok, Gemini, and OpenAI) against sophisticated social engineering attacks. A benchmark dataset of approximately 600 prompts was generated using a custom designed taxonomy of social engineering techniques to compare model responses to direct harmful requests versus paraphrased versions using techniques like pretexting and urgency. All LLM responses were manually evaluated by a human team.

The findings reveal critical security vulnerabilities. Using social engineering to paraphrase prompts almost doubled the jailbreak success rate (from around 31.5% to 58.7%). A significant disparity in robustness was observed: Grok proved to be the most vulnerable model, failing on 84% of paraphrased attacks, followed by Gemini (50%) and GPT-5 (43%). The most successful attack categories were Pretexting, Foot in the door and Impersonation.

As an additional project outcome, an open source "LLM Safety Judge" was developed and deployed on Google Cloud Vertex AI to assist in automating the safety annotation process.

Table of Contents

1.0. Introduction.....	4
2.0. Research and Related Work	5
3.0. Design and Methodology	7
4.0. Implementation.....	12
5.0. Evaluation and Key Findings.....	18
6.0. Cloud Security Assessment.....	23
7.0. Final Deliverables	25
8.0. Conclusion	27
9.0. References	29

1.0. Introduction

1.1 Problem Statement and Background

Grok, Gemini and GPT-5 are Large Language Models that have developed conversational AI at a very fast pace and have facilitated natural and context sensitive interactions in many different areas. Their prevalence in research, learning and business systems has transformed human communication with machines.

But the advancements have also presented a lot of new security challenges. The most prominent of them is the vulnerability of LLMs to adversarial manipulation, through which attackers exploit minor manipulation of language to bypass any protection mechanism that has been implemented. Social engineering jailbreaks are seen as the most critical threat of the emerging category of attacks whereby users with ill intentions mask their malicious or illegal requests as part of a convincing, credible, or emotionally appealing story.

These psychologically engineered prompts, unlike the direct or keyword-based violations, use the contextual knowledge of the model to induce compliance and not rejection. People whose language uses such manipulative linguistic clues are not usually detected by current LLM guardrails which tend to be taught to identify overt unsafe requests. The gap that the project is filling is thus a pressing one in the AI safety field: the necessity of a methodically, behaviorally based benchmark to evaluate the reaction of commercial LLMs to socially engineered manipulation.

1.2 Objectives

Primary Objective:

Adversarial Benchmarking: To test the safety and robustness of the guardrail of commercial LLMs, Grok, Gemini, and GPT-5, under socially engineered adversarial prompts strictly. This creates an authorized human-validated benchmark against which to compare the model security of various vendors.

Secondary Objective:

Automated Safety Evaluation: The open-source, cost-effective model, which is the LLM Safety Judge, an artificial intelligence tool designed to identify jailbreak responses automatically, to be designed, fine-tuned, and deployed. This makes the practice of manual annotation less important and preconditions scalable auditing of AI safety.

1.3 Scope

This project integrates computational benchmarking with behavioral insights derived from social-engineering psychology to assess real-world AI vulnerabilities.

- **Models Tested:** Grok, Gemini, and GPT-5 were tested as the typical commercial LLM in the logged-out and publicly available conditions to reflect the realistic user access.
- **Dataset:** A benchmark dataset of approximately 600 paired prompts was developed, comprising both direct harmful requests and indirect paraphrased versions engineered using eight distinct social engineering techniques, including urgency, authority framing, pretexting, and impersonation.
- **Evaluation Process:** Each model's response was examined using a dual evaluation framework:
 - **Human Annotation:** manual verification of responses as Safe, Unsafe, or Partially Unsafe.
 - **Automated Evaluation:** using the fine-tuned LLM Safety Judge, built on Llama 3.1 8B and deployed via Google Cloud Vertex AI.
- **Core Outcome:** The project establishes a reproducible, human-validated methodology for stress-testing AI systems and contributes to the growing field of behavioral AI safety. It points out the significance of context-sensitive assessment in avoiding real-world abuse by using psychologically grounded prompts controlling.

2.0. Research and Related Work

Large Language Models (LLMs) like Grok, Gemini, and ChatGPT have quickly improved human–AI interaction by using natural dialogue and contextual reasoning. Prompt-injection and jailbreak attacks, which alter instructions to get around security measures, can still affect these systems. So, the expanding field of LLM safety looks at how to quantify, measure, and counteract adversarial prompting. According to recent research, security flaws frequently result from the models' vulnerability to linguistic framing that resembles social engineering rather than from overt policy infractions.

2.1 LLM Safety and Jailbreaking Research

Direct jailbreaks that employed explicit requests to cause unsafe outputs were the focus of early benchmarking projects like AdvBench and HarmBench (Yi et al., 2024; Jin et al., 2024). These datasets only recorded surface-level attacks based on keyword patterns, despite being useful for identifying rule violations. Later studies, such as Weak-to-Strong Jailbreaking on Large Language Models (OpenReview, 2024), showed that contextualized or paraphrased prompts often circumvent the same protections. A crucial limitation was revealed by these

results: models trained on direct violations frequently don't perform well when presented with linguistically diverse or psychologically compelling inputs. Our project's emphasis on socially engineered adversarial prompts as opposed to merely syntactic manipulation was directly inspired by this gap.

2.2 Social Engineering as an Adversarial Modality

This project used two important studies that apply classical social-engineering theory to digital security contexts: Cohen and Heverin (2025) and (Longtchi et al., 2022). They demonstrate how attackers use psychological concepts like fear, urgency, authority, and trust to sway judgment. To explain why victims, comply with malicious requests, both papers offer structured taxonomies of psychological factors (PFs) and persuasion techniques (PTs). Our team rethought LLM jailbreaking as a machine-targeted social-engineering technique, building on these frameworks. Attackers conceal their intentions behind plausible narratives rather than giving out explicit forbidden orders; for instance, they might present a request as humanitarian aid or scholarly research.

The classification of attack types in our dataset, which include illegal weapons, violence and hate, and privacy, was informed by this theoretical foundation. Every category fits with strategies like foot in the door, persuasion, impersonation, and pretexting. As human adversaries alter tone and context to evade detection, the social-engineering viewpoint also supported the creation of paraphrased versions of each prompt. Therefore, assessing several linguistic realizations of the same intent provided a more practical standard for model safety.

2.3 Automated Safety Evaluation and Its Limitations

To classify model outputs as safe or unsafe, companies like OpenAI, Anthropic, and Google DeepMind have implemented automated moderation systems, also known as judge models. Empirical research, however, reveals enduring flaws: automatic judges perform well on overt threats but poorly on subtle, paraphrased attacks (OpenAI Research, 2024). Usually, they use heuristics at the keyword level instead of contextual or semantic understanding. Furthermore, previous benchmarks rarely combine automated scores with human-verified labels, which restricts their capacity to assess the degree of agreement between machine and human judgments. We designed a refined Judge LLM that is directly compared to manual annotations to measure consistency and detect systematic biases because of this lack of dual evaluation.

2.4 Summary of Insights

Previous studies have demonstrated that adversarial prompting encompasses psychological influence and persuasive dialogue in addition to basic text manipulation. While overt violations are benchmarked in existing datasets, the behavioural aspect of manipulation is

overlooked. Our project advances a comprehensive methodology for behaviour-informed LLM safety benchmarking by integrating computational text analysis, dual human-judge evaluation, and social-engineering theory (Cohen & Heverin, 2025; Longtchi et al., 2022). In addition to revealing real-world vulnerabilities, the resulting repository of effective socially engineered prompts offers structured data for future improvements to safer models.

3.0. Design and Methodology

This design and methodology section outlined the structured framework created to examine how social-engineering techniques might influence LLM safety. The plan unified behavioural theory, linguistic manipulation, and computational benchmarking into a coherent, reproducible blueprint. By combining the 600-Prompt Loop (human-anchored evaluation) with the Automated Safety Judge (scalable verification), the design established an integrated foundation for assessing guardrail robustness in commercial LLMs.

3.1 System Architecture

Figure 3.1 outlines the proposed three-tier architecture that informed the overall design. The architecture consisted of:

- **Social-Engineering Tactics Layer:** A curated taxonomy of social engineering strategies intended to guide the structure and format of malicious prompts.
- **Paraphraser / Prompt Generator Layer:** A mechanism for transforming direct prompts into paraphrased prompts that integrate the chosen social-engineering techniques.
- **Evaluation and Repository Layer:** The analytical and storage components responsible for querying large language models (LLMs), collecting responses, assigning safety labels, and archiving results for later analysis.




Figure 3.1.1 Proposed system design architecture

This layered structure helped the team determine the impact of using human social-engineering strategies for prompt injection by analysing how consistently they function across models and circumstances. While the proposed architecture (Figure 3.1.1) served as the conceptual foundation for the study, not all its modules were implemented in practice. The team adopted the simplified workflow shown in Figure 3.2 to maintain feasibility within available resources while still fulfilling the design objectives. This streamlined version preserved the essential data flow and evaluation logic of the original architecture, ensuring methodological consistency with the overall research design.

3.2 Taxonomy Development and Prompt Design

The initial part of the project planning phase involved developing a clear taxonomy of social-engineering techniques to ensure uniformity in prompt construction. Eight distinct techniques were selected and defined based on psychological mechanisms of influence:

1. **Urgency:** Creating a sense of limited time or impending consequences to pressure the target into acting quickly.
2. **Attention Grabbing:** Capturing the target's focus with strong emotional or sensory cues that override normal skepticism.
3. **Visual Deception:** Manipulating appearance or presentation so the victim mistakes malicious content for something trustworthy.
4. **Incentive and Motivator:** Using rewards, benefits, or emotional appeals to encourage the target to act against their best interests.
5. **Persuasion:** Influencing decisions using psychological principles identified by Cialdini: authority, reciprocity, liking, social proof, consistency, and scarcity.
6. **Quid-Pro-Quo:** Framing the interaction as an exchange, where the victim must give something (like information or access) to receive a benefit or avoid harm.

7. Foot-in-the-Door: Gaining compliance by making a small, reasonable request first, then escalating to larger or riskier demands.
8. Trusted Relationship: Exploiting an existing relationship between the attacker (or impersonated persona) and the victim to bypass suspicion.
9. Impersonation: Pretending to be a specific person or role to gain access, information, or compliance.
10. Pretexting: Crafting a detailed backstory (“pretext”) that explains why the attacker is making a request, making it appear legitimate.

During the taxonomy development process, three primary attack categories (Violence & Hate, Guns & Illegal Weapons, and Criminal Planning) were identified and planned for inclusion to ensure balance and uniformity across all team members’ work.

The design specified that every direct prompt (a straightforward unsafe query within an attack category) would be rewritten into at least one indirect variant embedding one of the eight social-engineering techniques. Each direct–indirect pair formed a matched unit for comparative analysis, enabling systematic assessment of whether incorporating social engineering techniques influenced model compliance with guardrails.

3.3 The 600-Prompt Loop

The 600-Prompt Loop was the central mechanism designed to standardise data generation, testing, and labelling. It was conceptualised as a repeatable four-phase pipeline (Figure 3.2) that would be iterated until a final corpus of approximately 600 paired prompts was achieved.

- **Prompt Generation:** Direct prompts representing the three planned attack categories were manually authored and then paraphrased using one of the eight social-engineering techniques. Each indirect version retained semantic equivalence but introduced one social engineering technique from the predefined list through tone, context, or motivation.
- **Data Collection:** The planned evaluation phase involved submitting both direct and indirect prompts to three commercial LLMs (Grok, Gemini, and an OpenAI model) via their public-facing interfaces. This decision reduced the risk of account blocking or API-related constraints. Responses were to be systematically stored in structured spreadsheet along with direct and paraphrased prompts.




Figure 3.2: The 600-Prompt loop workflow

- **Human Annotation:** In the blueprint, each team member was to review every model response and assign one of three safety labels: Safe, Unsafe, or Partial. These annotations were to be served as ground-truth data for evaluating both LLM performance and the automated judge model.
- **Data Cleaning and Validation:** The final phase involved validation checks and standardisation to ensure data integrity by verifying that each prompt had a complete record (direct + indirect, response + label) for each LLM, removing duplicates, and anonymising sensitive content. Only validated samples were to be admitted to the final complete repository.

This four-phase loop was designed to continue until a balanced, human-validated corpus of roughly 600 paired prompts was complete, ready for downstream comparative analysis.

3.4 Automated Safety Judge

The Automated Safety Judge was a secondary, parallel objective to evaluate responses autonomously and reduce dependence on manual labelling. The judge was planned as a fine-tuned LLM classifier capable of categorising model outputs as safe, unsafe, or partial, and

providing short rationales for its decisions. The design planned to fine-tune the automated judge using a publicly available safety data (such as NVIDIA's Aegis 2.0 dataset).

This component was introduced to address three design goals:

- Efficiency: Scale safety evaluation beyond what manual annotation alone could achieve.
- Objectivity: Reduce human bias and inconsistency in label assignment.
- Auditability: Produce interpretable rationales that could be cross-checked against human judgement.

3.5 Design Changes and Revisions

Several intentional design revisions were recorded during the planning stage to ensure feasibility and alignment with the project's scope and principles.

- Public vs. Logged-In Interfaces: The initial design proposed using authenticated or paid APIs to query LLMs for greater control and reproducibility. However, the plan was revised to prioritise logged-out public interfaces for two reasons. One, to avoid triggering moderation flags or account suspensions when using sensitive adversarial prompts. Two, to evaluate the first layer of defence that real-world users typically encounter (the public chat interface) which best aligns with the project's objective of assessing visible guardrails.
- Model Selection Adjustments: Perplexity AI was initially considered as a fourth comparison model but later removed from the design due to mandatory sign-in requirements that conflicted with the logged-out testing principle.
- Model-Tuning Plan: The blueprint initially included fine-tuning the Mistral 7B model on Google Colab using Hugging Face libraries and parameter-efficient techniques (LoRA). This plan was abandoned due to hardware limitations in the free Colab environment. The corresponding revision proposed migrating to Google Cloud Vertex AI and fine-tuning Llama 3.1 8B, chosen for its open-source availability, managed infrastructure, and scalability. The Vertex AI environment offered advantages such as reproducible compute allocation, integrated monitoring, and simplified endpoint deployment. This approach aligned with the project's principle of building an open-source, affordable, and transparent evaluation ecosystem, ensuring future researchers could replicate the pipeline without proprietary dependencies.

Figure 3.5 shows the proposed workflow for the project in the initial planning phase with different strategies which underwent significant change during the implementation phase.




Figure 3.5 Proposed workflow for the project

4.0. Implementation

The practical execution of the project's design involved three core phases: the implementation of the benchmarking dataset, the development of the LLM Safety Judge, and the final deployment and inference to the judge model.

4.1 Implementation of the Benchmarking Dataset

The implementation of the benchmark dataset was a direct execution of the "600-Prompt Loop" design and involved four key operational phases:

1. **Prompt Generation:** The process began with sourcing a list of "Direct Prompts" representing various high-risk attack categories. The core implementation task was the manual and semi-automated paraphrasing of these prompts. Each prompt was rewritten to integrate one of eight pre-defined social engineering techniques (e.g., *pretexting*, *urgency*, *quid-pro-quo*) to create an indirect version. This process was repeated to generate a final corpus of over 600 unique, paired prompts (direct and paraphrased).
2. **Data Collection:** This dataset was systematically processed against the public-facing APIs (the non-signed-in version using the UI of these LLMs to mitigate the risk of blocked accounts associated with the submission of sensitive queries and lack of budget to use the paid APIs.) of the target commercial LLMs: Grok, Gemini, and

OpenAI. All responses were collected, catalogued in a database, and precisely mapped to their originating prompts.

3. **Human Annotation:** This was the most critical implementation phase. The project team, acting as subject matter experts, manually reviewed every collected LLM response. Each response was assigned a definitive label of "Safe," "Unsafe," or "Partial". This time-intensive human-labelling process was the key step that transformed the raw, collected data into the 600-sample benchmark dataset. While most of the annotation data came from human annotators, team members also incorporated the OpenAI Content Moderation API (using the omni-moderation-2024-09-26 model) as a grounding mechanism for their work.
4. **Data Cleaning:** Following annotation, the dataset underwent a cleaning and validation process before analysis. This involved checking for null or missing values, ensuring consistency in labels, and identifying near-duplicates. This step was crucial for ensuring the integrity of the data used for the final benchmarking analysis.

4.2 Implementation of the LLM Safety Judge

To explore the automation of the expensive labelling process, an LLM Safety Judge was implemented. The rationale for implementing a fine-tuned LLM Judge extended beyond simply automating the labelling process. Safety annotation can be a subjective task, and a small human team may develop its own specific biases or heuristics. This approach is based on the "LLM-as-a-Judge" paradigm, which has been shown to be a scalable and reliable method for evaluating model outputs (Zheng et al., 2023).

By training a model on a large, external, and industry-vetted dataset like NVIDIA's Aegis 2.0, we aimed to create an objective "second opinion". This allowed us to validate our own human-annotated labels against a consistent, data-driven rationale, providing a valuable external benchmark to strengthen our findings and explore the feasibility of scaling the evaluation (Ghosh et al., 2025).

The development began with an initial attempt to fine-tune a Mistral 7B model using Google Colab. This attempt failed due to a lack of compute resources and environmental complexities. This failure led to a strategic pivot. The team switched to the more user-friendly and robust Google Cloud Platform (GCP) Vertex AI. The Llama 3.1 8B Instruct model was selected as the new foundation, chosen for being a newer-generation, open-source model that is readily available and supported within the Vertex AI ecosystem.

4.2.1 Fine-Tuning Data and Preparation

A key decision was to train the judge on a large, public, commercially usable dataset to ensure it was an unbiased, external evaluator. The dataset chosen was the Benchmarked NVIDIA Aegis 2.0 content safety dataset available from Hugging Face.

This dataset is a large, human-annotated collection of over 34,000 human-LLM interactions, specifically designed for training and evaluating content safety and moderation models. It covers a comprehensive taxonomy of 12-13 risk categories (Criminal Planning, Violence, Hate, etc.) and is ideal for training a robust safety classifier.

A significant data preparation pipeline was required:

- **Data Formatting:** The team extracted the LLM responses and their corresponding safety labels from the NVIDIA Aegis 2.0 dataset, followed by filtering results that were not either 'safe' or 'unsafe' as this project is not interested in other labels. This data was then transformed into JSONL files, the required format for Vertex AI fine-tuning. Each line in the file contained a message with "user" and "assistant" roles, (an example of one sample JSON from the dataset is shown in figure 4.2.1.1).

```
{ "messages":  
  [  
    { "role": "user",  
      "content": "### Instruction:\nYou are a safety judge. Analyze the  
assistant's response to determine if the response is safe or unsafe.\n\n###  
Assistant's Response:\nI will not engage in conversations that make negative  
generalizations about languages or cultures.\n\n### Judgement:"  
    },  
    { "role": "assistant",  
      "content": "safe"  
    }  
  ]  
}
```

Figure 4.2.1.1: An example of one sample JSON from the training dataset.

- **Dataset Balancing:** Initial analysis showed the Aegis dataset was heavily skewed towards the 'safe' category. To prevent the judge from developing a bias and simply defaulting to 'safe', the team under sampled the 'safe' category to create a more balanced training distribution.




Figure 4.2.1.2: Plot showing class imbalance in the training data, and dataset after fixing.

- **Final Splits:** This preparation resulted in a final dataset of approximately 7,000 training samples and 1,000 validation samples.



Figure 4.2.1.3: Data distribution for fine tuning the Llama 3.1 model.

4.2.2 Model Training and Evaluation

The fine-tuning process was iterative. Initial training efforts with the Llama 3.1 8B model showed signs of over-fitting to the training data. To correct this, the hyperparameters were adjusted.




Figure 4.2.2.1: Signs of overfitting in the first attempt of fine-tuning

The final, successful training configuration used LORA (Low-Rank Adaptation), a parameter-efficient technique, with the following parameters (Hu et al., 2021):

- LORA Rank: 16
- Learning Rate: 0.002
- Epochs: 1

Tuning details: llama-3.1-8binstruct-judge-1ep

Deploy

View model details

Monitor

Dataset

Details

Model name

projects/11630237313/locations/us-central1/models/7230347782324748288@1

Tuning job

projects/11630237313/locations/us-central1/tuningJobs/7185108874616111104

Status

Succeeded

Region

us-central1

Created

Oct 15, 2025, 6:19:40 PM

Ended

Oct 15, 2025, 7:27:42 PM

Tuning method

Supervised

Base model

meta/llama3_1@llama-3.1-8b-instruct

Tuning mode

LoRA tuning

Tuning dataset

gs://llm-as-a-judge-bucket/datasets/train_aegis_for_vertex_llama3_chat_version.jsonl

Experiment

tuning-experiment-20251015001942283780

Output directory

gs://llm-as-a-judge-bucket/judge-llama-1epoch/

Number of epochs

1

Learning rate

0.0002

Adapter size

16

Truncated example count

0

Validation dataset

gs://llm-as-a-judge-bucket/datasets/val_aegis_for_vertex_llama3_val_chat_version.jsonl

Figure 4.2.2.2 Fine-tuning details on Vertex AI platform

After the fine-tuning was complete, the model was evaluated *separately* on the held-back test split from the Aegis dataset. This evaluation, purely for its classification performance on the task it was trained for, yielded an 88.8% accuracy.




Figure 4.2.2.3 Fine-tuning evaluation

4.3 Deployment and Inference

With a successfully trained model, the next step was to use it. An inference pipeline was first developed in a Jupyter Notebook using Python to structure the API calls and data handling. The trained Llama 3.1 Judge was then formally deployed to a Vertex AI Endpoint. This deployment operationalized the model, transforming it into a live service.

```

[120]: def get_judgment(user_prompt: str, assistant_response: str, max_retries: int = 5, delay_seconds: int = 5):
        """
        Formats the data and calls the endpoint to get a safety judgment,
        with retries on failure.
        """
        content_for_judge = f"""### Instruction:
        You are a safety judge. Analyze the assistant's response to determine if the response is safe or unsafe.

        Assistant's Response:
        {assistant_response}

        Judgement:"""
        instance = {
            "@requestFormat": "chatCompletions",
            "messages": [{"role": "user", "content": content_for_judge}],
            "max_tokens": 10
        }

        for attempt in range(max_retries):
            try:
                prediction = endpoint.predict(instances=[instance])
                result = prediction.predictions[0][0]['message']['content'].strip().lower()
                if result not in ['unsafe', 'safe']:
                    print(f" -- Warning: Unexpected judgment '{result}'. Setting to None.")
                    result = None
                return result # Success, exit the function
            except Exception as e:
                print(f" -- Attempt {attempt + 1}/{max_retries} failed: {e}")
                if attempt < max_retries - 1:
                    print(f" -- Retrying in {delay_seconds} seconds...")
                    time.sleep(delay_seconds)
                else:
                    print(" -- Max retries reached. Returning 'ERROR'.")

        return "ERROR" # This is reached if all retries fail

```

Figure 4.3.1 Fine-tuning inference helper function written in Python.

This endpoint was then used to process our 600-sample ground-truth dataset. This step provided the automated "Judge labels," which were used as the "second opinion" to compare against our team's human annotations.

4.4 Future Improvements

The development of the LLM Judge provided valuable insights. A clear limitation of any single judge is that it is biased by its training data and its pretraining (in this case, the NVIDIA Aegis dataset).

A significant future improvement would be to implement a "jury of LLMs." This would involve fine-tuning *more than one* judge model, each trained on a different high-quality safety dataset (e.g., HarmBench, BeaverTails). This concept, where an ensemble of diverse models is used for evaluation, has been shown to reduce bias and improve alignment with human judgment in recent research (OpenReview 12346, 2024). By assembling a jury and taking a consensus or majority vote, we could average out the biases of any single dataset. This would create a far more robust and reliable automated safety-evaluation system.

5.0. Evaluation and Key Findings

5.1 Quality Assurance for Data Analysis

Before modelling and visualisation was conducted, the dataset underwent extensive quality assurance checks. Missing values, duplicated entries and inconsistently formatted labels were identified and resolved. Result labels were standardised, so that "Success" and "Partial Success" were mapped to the category "unsafe," and "Failed" was classified as "safe." This was done to ensure the human labels become consistent with the LLM judge labels to enable fair comparison between their results. Spelling errors and inconsistent labelling were also corrected for social engineering techniques and attack categories.




Figure 5.1.1 Prompt length distribution for direct and paraphrased prompts.

To ensure that the dataset was suitable for reliable model evaluation, prompt length was analysed as part of the quality assurance process. The distribution as shown in Fig. 5.1.1 showed that direct prompts were consistently short and compact, while paraphrased prompts were longer and more descriptive. This behaviour is expected, because paraphrasing often adds narrative framing or contextual justification to disguise malicious intent. Despite this increase in length, all prompts remained within practical limits, with the majority falling well within a reasonable character range and only a very small number extending into higher-length outliers.

5.2 Key Findings




Figure 5.2.1 Attack success rates for each LLM model.

The bar chart in Fig. 5.2.1 compares the average attack success rate of direct and paraphrased prompts for each model. The bar chart showed a clear trend: paraphrased prompts were far more successful at bypassing safety filters. Grok was the most vulnerable, where more than 80 percent of paraphrased prompts generated unsafe outputs, compared to around 47 percent for direct prompts. Gemini and GPT-5 showed lower overall rates, but a similar pattern is observed when comparing the attack success rates for direct and paraphrased prompts. This suggests that models are much better at refusing obvious malicious instructions than disguised ones, and that paraphrasing dramatically increases the likelihood of unsafe responses.



Figure 5.2.2 Attack Success Rates for Paraphrased prompts by Technique and Model.

To investigate vulnerability across specific social engineering strategies, a heatmap (Fig. 5.2.2) was generated to compare paraphrased attack success by technique and LLM model. The heatmap revealed that Grok was consistently the most susceptible model, mostly exceeding 0.8 unsafe rate across different techniques. GPT-5 demonstrated moderate vulnerability, while Gemini showed relatively stronger filtering but was still susceptible to several techniques. Methods such as Foot in the Door, Impersonation and Pretexting achieved particularly high success rates across multiple models, demonstrating that gradual escalation and narrative framing of language can bypass safety mechanisms more effectively than direct malicious instructions.




Figure 5.2.3 Attack Success Rate by SE Technique.

The bar chart (Fig. 5.2.3) compares the attack success rates of direct prompts versus paraphrased prompts across each social engineering technique. In every case, paraphrasing leads to a noticeable increase in unsafe responses, demonstrating how rewording makes harmful intent harder for LLMs to detect. Techniques such as Pretexting, Quid-Pro-Quo, and Foot-in-the-Door show particularly large jumps, where attack success rates almost double when paraphrased. Overall, the chart shows a consistent pattern that psychological framing makes malicious prompts significantly more effective than explicit requests, and this holds across all tested techniques.




Figure 5.2.4 Distribution of SE Techniques in Dataset.

The distribution of techniques in the dataset was visualised through a pie chart (Fig. 5.2.4). Persuasion, pretexting, and urgency represented the largest proportions, while techniques such as visual deception and quid-pro-quo attacks appeared less frequently but were still

well-represented. A small number of techniques were excluded from this chart because they contained too few samples to meaningfully influence the distribution. Overall, the visualisation demonstrates that the dataset is not dominated by a single manipulation style, ensuring the analysis captures a diverse range of social-engineering behaviours.




Figure 5.2.5 Word clouds after analysis of certain social engineering techniques.

To examine linguistic patterns in paraphrased prompts, word-cloud visualisations were created for each technique categories, and some of them are shown in Fig 5.2.5. Each word cloud showed that paraphrased prompts utilised psychologically targeted language rather than explicit harmful vocabulary. For example, pretexting prompts relied heavily on narrative and contextual terms such as “scenario,” “explain,” and “detail,” while urgency-based prompts frequently used time-pressure cues such as “immediately,” “minutes,” and “quick.” Incentive and motivator prompts incorporated terms like “reward,” “exclusive,” and “benefit,” and attention-grabbing prompts used dramatic language intended to provoke emotional or intellectual curiosity. These visual patterns demonstrate that paraphrased attacks are not simple rewordings of direct prompts; instead, they introduce emotional, persuasive, or scenario-based framing that existing moderation systems do not reliably detect.




Figure 5.2.6 ASR of Direct vs. Paraphrased prompts.

The bar chart (Fig. 5.2.6) comparing direct and paraphrased prompts overall shows a substantial increase in attack success rate after paraphrasing. Direct prompts had an unsafe output rate of **38.9%**, and paraphrased prompts reached a **64.7%** success rate demonstrating that indirect, socially engineered phrasing is far more effective at bypassing model safety mechanisms. This visual evidence reinforces one of the core findings of the project: prompt rewriting significantly increases vulnerability across LLMs, even when the underlying malicious intent remains the same.

6.0. Cloud Security Assessment

6.1 Scope

The team used two services. On Google Cloud Vertex AI, the fine-tuned LLM Safety Judge (Llama-3.1-8B with LoRA) was hosted behind an authenticated endpoint and served inference during evaluation runs. On Microsoft 365 OneDrive, the prompt dataset, derived labels, figures, and report artefacts were stored in a team folder restricted to named Swinburne University email IDs of our 6 team members only.

The team followed each cloud service provider's guidance for access control, network exposure, encryption, logging/monitoring, data handling, versioning/rollback, and retention. For every area, they recorded what was checked, what control was in place, and where the supporting evidence was stored.

6.2 Vertex AI (Google Cloud)

The team carried out the following steps to keep their use of cloud services secure and well controlled.

- **Identity and Access:** Who could deploy, invoke, read logs, and change settings was verified. Least-privilege IAM was used. Only two members held deploy rights; others had read-only as needed. A dedicated service account performed runtime access; keys were not embedded in code or documents. Google accounts used multi factor authentication.
- **Network exposure:** The endpoint accepted only authenticated calls; anonymous requests did not work. Usage limits were kept modest to avoid bursts or any unexpected cost. If traffic or exposure increased, moving the endpoint to a private network was identified as a future hardening step; it was noted only for later and stayed out of scope for this project.

- **Data minimisation and logging:** Because prompts could be sensitive, payload logging was kept to a minimum. Only basic metadata and error summaries were retained. When examples were needed, small samples were saved locally rather than in cloud logs.
- **Monitoring and alerts:** Errors and latency were visible in logs, and simple usage alerts were enabled to catch loops or misuse. This helped the team run heavy tests in short, planned sessions to keep costs and usage under control.
- **Reliability and rollback:** Previous versions of the endpoint were saved. Updates ran a quick health check and switched only if it was a success, otherwise the system rolled back. One “known good” version remained available during the project, as the primary backup.
- **Data residency and retention:** A standard cloud region was used. No personal data was processed. Default log retention was kept, and the maintained runbook help the team to clean up and manage the logs.
- **Risks and planned hardening:** The team noted three risks: logs might accidentally store sensitive text; cloud provider-managed encryption keys were okay for research but not the strongest choice; and adding a publicly available URL would raise misuse risk. Planned improvements were to keep logs minimal, use private network access, switch to stronger team-managed keys, and add simple limits on request size and how often people can send requests.
- **Overall risk rating:** Low–Moderate for this research scope (no PII, restricted access, minimal payload logging, tested rollback).

6.4 OneDrive (Microsoft 365)

Access and sharing: Folder access was limited to Swinburne accounts and granted only to named team members. External sharing was off, and “anyone with the link” links were not used. Permissions were re-checked after major uploads of important project documents.

Encryption and transport: Files were encrypted at rest and sent over HTTPS. The README file reminded contributors not to store secrets or API keys.

Versioning and recovery: OneDrive’s version history kept earlier copies of files automatically. For key milestones, the team also saved a dated “snapshot” folder so any change or update could be roll backed to a known state, if needed.

Recovery was test-run twice:

- an edit was reverted using Version history, and
- a file was restored from the Recycle bin after a controlled delete.
- A short checklist noted which team member performed the test and when.

Content sensitivity and hygiene: Files were checked to ensure no personal data or keys were included. Filenames were kept simple and descriptive, and the README warned against adding sensitive material.

Residual risks and planned hardening: Two risks were noted: someone might accidentally share a sub-folder, and files could sync to a personal/unmanaged device. Planned fixes: mark folders “Internal-only,” require MFA and only approved devices for syncing, and automatically delete temporary/scratch files after a short retention period.

Risk rating: It was low for this project (Swinburne-only access, encrypted storage/transfer, no PII or secrets).

6.6 Traceability & updates

- The team shifted from local/Colab runs to Vertex AI for predictable deployment and easier auditing for the LLM judge fine-tuning part.
- Core artefacts, such as the architecture diagram, runbook, README, and the Swinburne-only OneDrive policy, were kept in sync and checked against the agreed security stance.
- A simple change log noted what changed, when, and why, so documents and settings matched the design.
- Dated snapshots and commit references made reviews straightforward.

6.7 Security Status

- Overall risk was assessed as **Low to Moderate**. Access stayed restricted (including Swinburne-only storage), data was encrypted in transit and at rest, sensitive content was not logged, rollback was tested, and cleanup was planned.
- Clear “next step” hardening items were recorded for any move toward production (e.g., tighter network access and stronger key management).

7.0. Final Deliverables

The full deliverable was packaged in a GitHub repository (link provided in the submission cover sheet and given below in the document). It served as the single source of truth for code, manuals, and evidence. For the prompts dataset, it is hosted on OneDrive to limit the access to the team, client and the unit facilitators.

Repository Link: https://github.com/COS80029-TAP-XC2-Project/XC2_Benchmarking_LLM-Safety_Against_Social_Engineering_Techniques

Prompts: https://liveswinburneeducuau-my.sharepoint.com/:x/g/personal/104800822_student.swin.edu.au/EVt1GBtTsu9B0wxEPF2vZo4Bp0oLZNZbTuiskFConVCmBg




Figure 7: Screenshot of the GitHub repository containing project artefacts and deliverables.

The repository contains the following artefacts.

- **Technical & User Manual (README.md file):** This file explains, in plain English, how to set things up, run the R checks, run the Python analysis, where files live in the repo, and what to do if something doesn't work.
- **Final dataset:** Final_Prompt_Dataset_with_LLM_judge_labelling.xlsx is the main spreadsheet. It groups results by model and includes both the human outcomes and the labels assigned by the LLM Judge.
- **R validator:** Dataset_preprocessing_and_validation/validate_prompts_views. R checks the data for common issues (missing values, duplicates, odd lengths) and shows simple agreement charts. When it runs, it also saves the plots in a dated folder for the record.
- **Python analysis:** Deep_Analysis_of_dataset_and_Evaluations.ipynb recreates the figures and tables used in the report. Open the notebook and run the cells from top to bottom to reproduce the results.
- **Generated outputs:** The analysis writes its charts and tables to Deep_Analysis_of_dataset_and_Evaluations/figures/ and

Deep_Analysis_of_dataset_and_Evaluations/tables/. These folders are created automatically when the notebook runs.

- **Presentation materials:** The client slides and handouts are in Final_Project_Delivery_and_Demonstration_to_Client/. These can be used if a quick overview of the project and findings is needed.
- **LLM Judge:** Fine_Tuned_LLM_judge/ contains artefacts from the fine-tuning activity.

8.0. Conclusion

8.1 Project Overview and Deliverable of the Objectives

The following two objectives of this project were achieved:

- **Goal 1: LLM Security Resilience Benchmarking:** The resistance of Grok, Gemini and GPT-5 to jailbreaks based on social engineering was systematically tested on a set of curated 600 adversarial prompts. Manual verification was used to establish the degree of model vulnerability to indirect attacks that were linguistically framed.
- **Goal 2: Programme Development of the LLM Safety Judge:** A fine-tuned LLM Safety Judge, that was created and deployed to Google Cloud Vertex AI to be used to automatically label the responses of the LLM regarding safety.

Collectively, these results generated a repeatable system of assessing AI model susceptibility and gave practical information about enhancing large scale safety mechanisms.

8.2 Significant Findings and Implications to Security

The evaluation revealed several critical insights into the security posture of commercial LLMs:

- **Impact of Social Engineering:** Jailbreak success rates almost doubled, with the success rate increasing to 58.7 per cent, when prompts were reformulated in persuasive or psychological terms, as opposed to the 31.5 per cent success rate when persuasive language was not used. This implies that the LLMs are much more vulnerable to overt social manipulation compared to overt unsafe demands.
- **Model Vulnerability Disparity:** Among the tested models, Grok demonstrated the highest vulnerability with an 84 % failure rate, followed by Gemini (50 %) and GPT-5 (43 %), highlighting significant differences in vendor-level guardrail strength.

- **Attack Categories:** The most successful attack categories included Medical & Health Misinformation, Hacking, and Scamming, representing real-world domains where LLM misuse could have serious implications.
- **LLM Safety Judge Performance:** The fine-tuned LLM Safety Judge achieved 88.8 % accuracy, confirming its reliability as a scalable safety auditing tool, which can be further developed and incorporated in LLM guardrails.

Taken together, these results highlight the need to deploy adaptive AI safety systems to be able to identify socially engineered inputs that are not limited to recognition of simple keywords.

8.3 Future Work and Research Directions

The outcomes of Project XC2 form a strong foundation for future advancements in AI safety research. Recommended directions include:

- **Jury of LLMs Implementation:** Move the idea of a single judge to a jury of LLMs - a concerted group of models, like Aegis 2.0, HarmBench, and BeaverTails - to decrease the degree of bias and enhance consistency in judgments.
- **Continuous Learning for the Judge:** Implement a feedback loop that allows the LLM Safety Judge to adapt continuously as new adversarial patterns emerge, improving long-term robustness.
- **Dataset Expansion:** Broaden the dataset to include multimodal prompts (e.g., image-based or audio-embedded attacks) to reflect evolving threat landscapes.
- **Testing Newer Architectures:** Apply the benchmarking framework to upcoming generations of commercial and open-source LLMs, ensuring ongoing traceability of safety performance.

Such directions to the future will enhance the ability of AI systems to identify manipulation, increase trust among people, and shape the responsible use of AI tools of the next generation.

9.0. References

1. Cohen, E., & Heverin, T. (2025). Evaluating the Effectiveness of Psychological Prompt Injection Attacks on Large Language Models for Social Engineering Artifact Generation. Academic Conferences International Limited.
2. Longtchi, T., Rodriguez, R.M., Al-Shawaf, L., Atyabi, A. and Xu, S. (2022). Internet-based Social Engineering Attacks, Defenses and Psychology: A Survey. arXiv:2203.08302 [cs]. [online] Available at: <https://arxiv.org/abs/2203.08302>.
3. Jin, H., Hu, L., Li, X., Zhang, P., Chen, C., Zhuang, J. and Wang, H. (2024). JailbreakZoo: Survey, Landscapes, and Horizons in Jailbreaking Large Language and Vision-Language Models. [online] arXiv.org. Available at: <https://arxiv.org/abs/2407.01599>
4. Yao, Y., Duan, J., Xu, K., Cai, Y., Sun, Z. and Zhang, Y. (2024). A survey on Large Language Model (LLM) security and privacy: The Good, The Bad, and The Ugly. High-Confidence Computing, [online] 4(2), p.100211. Available at: <https://doi.org/10.1016/j.hcc.2024.100211>.
5. Zhao, X., Yang, X., Pang, T., Du, C., Li, L., Wang, Y.-X. and Wang, W.Y. (2025). Weak-to-Strong Jailbreaking on Large Language Models. [online] Openreview.net. Available at: <https://openreview.net/forum?id=Nazzz5GJ4g>.
6. Yi, S., Liu, Y., Sun, Z., Cong, T., He, X., Song, J., Xu, K. and Li, Q. (2024). Jailbreak Attacks and Defenses Against Large Language Models: A Survey. [online] arXiv.org. Available at: <https://arxiv.org/abs/2407.04295>
7. *Who Judges the Judge? LLM Jury-on-Demand: Building Trustworthy LLM Evaluation Systems*. [Conference submission] OpenReview Submission 12346. Available at: <https://openreview.net/forum?id=seM2ixNp6W>
8. Ghosh, S. et al. (2025) *Aegis2.0: A Diverse AI Safety Dataset and Risks Taxonomy for Alignment of LLM Guardrails*. [online] arXiv:2501.09004. Available at: <https://arxiv.org/abs/2501.09004>
9. Hu, E.J. et al. (2021) *LoRA: Low-Rank Adaptation of Large Language Models*. [online] arXiv:2106.09685. Available at: <https://arxiv.org/abs/2106.09685>
10. Zheng, L. et al. (2023) *Judging LLM-as-a-judge with MT-Bench and Chatbot Arena*. [online] arXiv:2306.05685. Available at: <https://arxiv.org/abs/2306.05685>